

In [2]:

```
import sys
import os
import json
import pandas
import numpy as np
import optparse

from keras.callbacks import TensorBoard
from keras import regularizers
from keras.models import Sequential, load_model
from keras.layers import LSTM, Dense, Dropout, Conv1D, GlobalMaxPooling1D, Flatten, MaxPooling1D
from keras.layers.embeddings import Embedding
from keras.preprocessing import sequence
from keras.preprocessing.text import Tokenizer
from keras.utils import to_categorical
from collections import OrderedDict

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV, RandomizedSearchCV
from sklearn.metrics import confusion_matrix, classification_report, roc_auc_score
from random import sample
import re
import matplotlib.pyplot as plt

import nltk
# nltk.download('stopwords')
from nltk.corpus import stopwords
```

In [3]:

```
def plot_history(history):
    acc = history.history['accuracy']
    val_acc = history.history['val_accuracy']
    loss = history.history['loss']
    val_loss = history.history['val_loss']
    x = range(1, len(acc) + 1)

    plt.figure(figsize=(12, 5))
    plt.subplot(1, 2, 1)
    plt.plot(x, acc, 'b', label='Training acc')
    plt.plot(x, val_acc, 'r', label='Validation acc')
    plt.title('Training and validation accuracy')
    plt.legend()
    plt.subplot(1, 2, 2)
    plt.plot(x, loss, 'b', label='Training loss')
    plt.plot(x, val_loss, 'r', label='Validation loss')
    plt.title('Training and validation loss')
    plt.legend()
```

In [33]:

```
# find the longest text to determine max_log_length
#length = [len(i) for i in X_encoded]
# max(length)
processed_data = pandas.read_csv("data_1w.csv")
processed_data
```

Out[33]:

	Unnamed: 0	0	1	2	3	4	5	6	7	8	...	2040	2041	2042	2043	2044	2045	2046	2047	Score	Label
0	0	0	0	0	0	0	0	0	0	0	...	191	213	64	343	213	64	3026	753	1	0
1	1	0	0	0	0	0	0	0	0	0	...	316	24	702	18	277	149	363	1674	1	0
2	2	0	0	0	0	0	0	0	0	0	...	696	223	106	532	495	16	1062	2018	1	0
3	3	0	0	0	0	0	0	0	0	0	...	18	59	199	287	298	1	726	484	1	0
4	4	0	0	0	0	0	0	0	0	0	...	3507	128	10	73	507	562	277	149	1	0
...
9995	9995	0	0	0	0	0	0	0	0	0	...	914	3	75	360	56	510	6778	14	4	1
9996	9996	0	0	0	0	0	0	0	0	0	...	170	2393	455	53	417	42	130	292	4	1
9997	9997	0	0	0	0	0	0	0	0	0	...	173	120	788	352	450	16	45	318	4	1
9998	9998	0	0	0	0	0	0	0	0	0	...	409	336	42	48	23	6	134	18865	4	1
9999	9999	0	0	0	0	0	0	0	0	0	...	275	693	50	34	627	782	418	514	4	1

10000 rows × 2051 columns

In [34]:

```
# pad the data as each observation has a different length
#max_log_length = 2048 # larger than 2003
#X_processed = sequence.pad_sequences(X_encoded, maxlen=max_log_length)
```

In [35]:

```
#X_processed
y = processed_data[["Label"]].values
avg_data = processed_data.iloc[:, 1:-2].values
# y -= 1
avg_data
```

Out[35]:

```
array([[ 0,  0,  0, ...,  64, 3026,  753],
       [ 0,  0,  0, ..., 149,  363, 1674],
       [ 0,  0,  0, ...,  16, 1062, 2018],
       ...,
       [ 0,  0,  0, ...,  16,   45,  318],
       [ 0,  0,  0, ...,   6,  134, 18865],
       [ 0,  0,  0, ..., 782,  418,  514]], dtype=int64)
```

In [36]:

```
avg_data = np.array(avg_data)
y = y.reshape((y.shape[0], 1))
```

In [37]:

```
# split 70% train and 30% test data
X_train, X_test, y_train, y_test = train_test_split(avg_data, y, test_size=0.30, random_state=0)
y_train_cat = to_categorical(y_train)
y_test_cat = to_categorical(y_test)
```

In []:

XGBoost - Binary - tokenizer

In [7]:

```
import xgboost as xgb
```

In [8]:

```
xg_rdsearch = xgb.XGBClassifier()
```

In [9]:

```
param_rdsearch_xg = {'n_estimators':[100, 150, 200, 250, 300, 350, 400, 450, 500,550, 600, 650, 700, 750,
                                     850, 900, 950, 1000],
                     'learning_rate':[0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1, 1.1, 1.2, 1.3, 1.4,
```

```

'max_depth':[1,2],
'gamma': [0, 0.25, 0.5, 0.75, 1, 1.25, 1.5, 1.75, 2, 2.25, 2.5, 2.75, 3, 3.25, 3.5, 3
4.5, 4.75, 5]}

```

+

```

xg_rdsearch = RandomizedSearchCV(xg_rdsearch, param_rdsearch_xg, cv = 5, scoring = 'accuracy',
                                refit = True, n_jobs=-1, verbose = 5)

```

In [10]:

```

xg_rdsearch.fit(X_train, y_train)

```

Fitting 5 folds for each of 10 candidates, totalling 50 fits

C:\Users\14264\Anaconda3\envs\tf_gpu\lib\site-packages\xgboost\sklearn.py:888: UserWarning: The use of label encoder in XGBClassifier is deprecated and will be removed in a future release. To remove this warning, do the following: 1) Pass option use_label_encoder=False when constructing XGBClassifier object; and 2) Encode your labels (y) as integers starting with 0, i.e. 0, 1, 2, ..., [num_class - 1].

warnings.warn(label_encoder_deprecation_msg, UserWarning)

C:\Users\14264\Anaconda3\envs\tf_gpu\lib\site-packages\sklearn\utils\validation.py:63:

DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

return f(*args, **kwargs)

[00:35:33] WARNING: ..\src\learner.cc:1061: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

Out[10]:

```

RandomizedSearchCV(cv=5,
                  estimator=XGBClassifier(base_score=None, booster=None,
                                          colsample_bylevel=None,
                                          colsample_bynode=None,
                                          colsample_bytree=None, gamma=None,
                                          gpu_id=None, importance_type='gain',
                                          interaction_constraints=None,
                                          learning_rate=None,
                                          max_delta_step=None, max_depth=None,
                                          min_child_weight=None, missing=nan,
                                          monotone_constraints=None,
                                          n_estimators=100,...
                                          verbosity=None),
                  n_jobs=-1,
                  param_distributions={'gamma': [0, 0.25, 0.5, 0.75, 1, 1.25,
                                                  1.5, 1.75, 2, 2.25, 2.5, 2.75,
                                                  3, 3.25, 3.5, 3.75, 4, 4.25,
                                                  4.5, 4.75, 5],
                                      'learning_rate': [0.1, 0.2, 0.3, 0.4,
                                                         0.5, 0.6, 0.7, 0.8,
                                                         0.9, 1, 1.1, 1.2, 1.3,
                                                         1.4, 1.5, 1.6],
                                      'max_depth': [1, 2],
                                      'n_estimators': [100, 150, 200, 250,
                                                         300, 350, 400, 450,
                                                         500, 550, 600, 650,
                                                         700, 750, 800, 850,
                                                         900, 950, 1000]},
                  scoring='accuracy', verbose=5)

```

In [11]:

```

xg_bestparams = xg_rdsearch.best_params_
xg_bestparams

```

Out[11]:

```

{'n_estimators': 500, 'max_depth': 2, 'learning_rate': 0.5, 'gamma': 0.75}

```

In [12]:

```

xg_bestmodel = xg_rdsearch.best_estimator_
xg_bestmodel

```

Out[12]:

```
XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
              colsample_bynode=1, colsample_bytree=1, gamma=0.75, gpu_id=-1,
              importance_type='gain', interaction_constraints='',
              learning_rate=0.5, max_delta_step=0, max_depth=2,
              min_child_weight=1, missing=nan, monotone_constraints='()',
              n_estimators=500, n_jobs=8, num_parallel_tree=1, random_state=0,
              reg_alpha=0, reg_lambda=1, scale_pos_weight=1, subsample=1,
              tree_method='exact', validate_parameters=1, verbosity=None)
```

In [13]:

```
y_train_pred = xg_bestmodel.predict(X_train)
```

In [14]:

```
print(classification_report(y_train_pred,y_train))
```

	precision	recall	f1-score	support
0	0.92	0.91	0.92	3503
1	0.91	0.92	0.92	3497
accuracy			0.92	7000
macro avg	0.92	0.92	0.92	7000
weighted avg	0.92	0.92	0.92	7000

In [15]:

```
yhat_xg_bestmodel = xg_bestmodel.predict(X_test)
```

In [16]:

```
classi_reprt_xg = classification_report(y_test, yhat_xg_bestmodel)
```

In [17]:

```
print(classi_reprt_xg)
```

	precision	recall	f1-score	support
0	0.76	0.75	0.75	1511
1	0.75	0.75	0.75	1489
accuracy			0.75	3000
macro avg	0.75	0.75	0.75	3000
weighted avg	0.75	0.75	0.75	3000

CNN - Binary - Tokenizer

In [52]:

```
CNN_Model = Sequential()
CNN_Model.add(Conv1D(10, 20, activation = 'sigmoid', input_shape = (2048, 1)))
CNN_Model.add(Conv1D(20, 8, activation = 'sigmoid'))
CNN_Model.add(Conv1D(30, 5, activation = 'sigmoid'))
CNN_Model.add(GlobalMaxPooling1D())
CNN_Model.add(Dense(10, activation = 'sigmoid'))
CNN_Model.add(Dense(1, activation = 'sigmoid'))
CNN_Model.compile(optimizer = 'adam', loss = 'binary_crossentropy', metrics = ['accuracy'])
CNN_Model.summary()
```

Model: "sequential_10"

Layer (type)	Output Shape	Param #
conv1d_22 (Conv1D)	(None, 2029, 10)	210
conv1d_23 (Conv1D)	(None, 2022, 20)	1620
conv1d_24 (Conv1D)	(None, 2018, 30)	3030
global_max_pooling1d_4 (Glob	(None, 30)	0
dense_13 (Dense)	(None, 10)	310
dense_14 (Dense)	(None, 1)	11
Total params: 5,181		
Trainable params: 5,181		
Non-trainable params: 0		

In [53]:

```
CNN_history = CNN_Model.fit(X_train.reshape((X_train.shape[0], X_train.shape[1], 1)), y_train, epochs = 1
                             (X_test.reshape((X_test.shape[0], X_test.shape[1], 1)), y_test))
```

Train on 7000 samples, validate on 3000 samples

Epoch 1/100

7000/7000 [=====] - 1s 185us/step - loss: 0.6936 - accuracy: 0.5013 - val_loss: 0.6932 - val_accuracy: 0.5037

Epoch 2/100

7000/7000 [=====] - 1s 142us/step - loss: 0.6935 - accuracy: 0.5010 - val_loss: 0.6928 - val_accuracy: 0.5010

Epoch 3/100

7000/7000 [=====] - 1s 142us/step - loss: 0.6932 - accuracy: 0.5029 - val_loss: 0.6927 - val_accuracy: 0.5093

Epoch 4/100

7000/7000 [=====] - 1s 143us/step - loss: 0.6931 - accuracy: 0.5101 - val_loss: 0.6924 - val_accuracy: 0.5010

Epoch 5/100

7000/7000 [=====] - 1s 145us/step - loss: 0.6920 - accuracy: 0.5313 - val_loss: 0.6919 - val_accuracy: 0.5490

Epoch 6/100

7000/7000 [=====] - 1s 152us/step - loss: 0.6917 - accuracy: 0.5300 - val_loss: 0.6915 - val_accuracy: 0.5437

Epoch 7/100

7000/7000 [=====] - 1s 132us/step - loss: 0.6906 - accuracy: 0.5244 - val_loss: 0.6907 - val_accuracy: 0.5373

Epoch 8/100

7000/7000 [=====] - 1s 132us/step - loss: 0.6888 - accuracy: 0.5481 - val_loss: 0.6900 - val_accuracy: 0.5287

Epoch 9/100

7000/7000 [=====] - 1s 132us/step - loss: 0.6897 - accuracy: 0.5354 - val_loss: 0.6905 - val_accuracy: 0.5413

Epoch 10/100

7000/7000 [=====] - 1s 132us/step - loss: 0.6877 - accuracy: 0.5493 - val_loss: 0.6889 - val_accuracy: 0.5473

Epoch 11/100

7000/7000 [=====] - 1s 132us/step - loss: 0.6866 - accuracy: 0.5474 - val_loss: 0.6888 - val_accuracy: 0.5360

Epoch 12/100

7000/7000 [=====] - 1s 134us/step - loss: 0.6846 - accuracy: 0.5574 - val_loss: 0.6882 - val_accuracy: 0.5373

Epoch 13/100

7000/7000 [=====] - 1s 137us/step - loss: 0.6840 - accuracy: 0.5543 - val_loss: 0.6929 - val_accuracy: 0.5330

Epoch 14/100

7000/7000 [=====] - 1s 133us/step - loss: 0.6839 - accuracy: 0.5599 - val_loss: 0.6921 - val_accuracy: 0.5400

Epoch 15/100

7000/7000 [=====] - 1s 136us/step - loss: 0.6839 - accuracy: 0.5569 - val_loss: 0.6869 - val_accuracy: 0.5483

Epoch 16/100

7000/7000 [=====] - 1s 138us/step - loss: 0.6818 - accuracy: 0.5706 - val_loss: 0.6861 - val_accuracy: 0.5477

Epoch 17/100

7000/7000 [=====] - 1s 133us/step - loss: 0.6819 - accuracy: 0.5689 - val_loss: 0.6860 - val_accuracy: 0.5497

Epoch 18/100
7000/7000 [=====] - 1s 138us/step - loss: 0.6811 - accuracy: 0.5663 - val_loss: 0.6867 - val_accuracy: 0.5427
Epoch 19/100
7000/7000 [=====] - 1s 181us/step - loss: 0.6801 - accuracy: 0.5687 - val_loss: 0.6869 - val_accuracy: 0.5457
Epoch 20/100
7000/7000 [=====] - 1s 146us/step - loss: 0.6808 - accuracy: 0.5653 - val_loss: 0.6878 - val_accuracy: 0.5430
Epoch 21/100
7000/7000 [=====] - 1s 133us/step - loss: 0.6806 - accuracy: 0.5731 - val_loss: 0.6906 - val_accuracy: 0.5457
Epoch 22/100
7000/7000 [=====] - 1s 132us/step - loss: 0.6810 - accuracy: 0.5627 - val_loss: 0.6883 - val_accuracy: 0.5450
Epoch 23/100
7000/7000 [=====] - 1s 156us/step - loss: 0.6805 - accuracy: 0.5656 - val_loss: 0.6857 - val_accuracy: 0.5467
Epoch 24/100
7000/7000 [=====] - 1s 143us/step - loss: 0.6793 - accuracy: 0.5657 - val_loss: 0.6844 - val_accuracy: 0.5513
Epoch 25/100
7000/7000 [=====] - 1s 146us/step - loss: 0.6789 - accuracy: 0.5691 - val_loss: 0.6846 - val_accuracy: 0.5510
Epoch 26/100
7000/7000 [=====] - 1s 138us/step - loss: 0.6787 - accuracy: 0.5681 - val_loss: 0.6852 - val_accuracy: 0.5513
Epoch 27/100
7000/7000 [=====] - 1s 133us/step - loss: 0.6796 - accuracy: 0.5680 - val_loss: 0.6852 - val_accuracy: 0.5583
Epoch 28/100
7000/7000 [=====] - 1s 132us/step - loss: 0.6793 - accuracy: 0.5700 - val_loss: 0.6901 - val_accuracy: 0.5457
Epoch 29/100
7000/7000 [=====] - 1s 132us/step - loss: 0.6778 - accuracy: 0.5761 - val_loss: 0.6846 - val_accuracy: 0.5573
Epoch 30/100
7000/7000 [=====] - 1s 132us/step - loss: 0.6771 - accuracy: 0.5770 - val_loss: 0.6850 - val_accuracy: 0.5483
Epoch 31/100
7000/7000 [=====] - 1s 132us/step - loss: 0.6769 - accuracy: 0.5739 - val_loss: 0.6854 - val_accuracy: 0.5480
Epoch 32/100
7000/7000 [=====] - 1s 132us/step - loss: 0.6770 - accuracy: 0.5759 - val_loss: 0.6852 - val_accuracy: 0.5560
Epoch 33/100
7000/7000 [=====] - 1s 132us/step - loss: 0.6763 - accuracy: 0.5790 - val_loss: 0.6852 - val_accuracy: 0.5640
Epoch 34/100
7000/7000 [=====] - 1s 132us/step - loss: 0.6774 - accuracy: 0.5751 - val_loss: 0.6842 - val_accuracy: 0.5537
Epoch 35/100
7000/7000 [=====] - 1s 132us/step - loss: 0.6770 - accuracy: 0.5773 - val_loss: 0.6856 - val_accuracy: 0.5540
Epoch 36/100
7000/7000 [=====] - 1s 132us/step - loss: 0.6754 - accuracy: 0.5837 - val_loss: 0.6840 - val_accuracy: 0.5593
Epoch 37/100
7000/7000 [=====] - 1s 137us/step - loss: 0.6753 - accuracy: 0.5816 - val_loss: 0.6843 - val_accuracy: 0.5573
Epoch 38/100
7000/7000 [=====] - 1s 132us/step - loss: 0.6770 - accuracy: 0.5760 - val_loss: 0.6845 - val_accuracy: 0.5610
Epoch 39/100
7000/7000 [=====] - 1s 132us/step - loss: 0.6760 - accuracy: 0.5800 - val_loss: 0.6857 - val_accuracy: 0.5540
Epoch 40/100
7000/7000 [=====] - 1s 137us/step - loss: 0.6757 - accuracy: 0.5794 - val_loss: 0.6897 - val_accuracy: 0.5433
Epoch 41/100
7000/7000 [=====] - 1s 142us/step - loss: 0.6760 - accuracy: 0.5764 - val_loss: 0.6868 - val_accuracy: 0.5483
Epoch 42/100
7000/7000 [=====] - 1s 137us/step - loss: 0.6763 - accuracy: 0.5756 - val_loss: 0.6845 - val_accuracy: 0.5500
Epoch 43/100
7000/7000 [=====] - 1s 149us/step - loss: 0.6759 - accuracy: 0.5784 - val_loss:

0.6860 - val_accuracy: 0.5560
Epoch 44/100
7000/7000 [=====] - 1s 135us/step - loss: 0.6745 - accuracy: 0.5804 - val_loss: 0.6856 - val_accuracy: 0.5543
Epoch 45/100
7000/7000 [=====] - 1s 137us/step - loss: 0.6764 - accuracy: 0.5809 - val_loss: 0.6858 - val_accuracy: 0.5557
Epoch 46/100
7000/7000 [=====] - 1s 139us/step - loss: 0.6747 - accuracy: 0.5820 - val_loss: 0.6861 - val_accuracy: 0.5540
Epoch 47/100
7000/7000 [=====] - 1s 151us/step - loss: 0.6749 - accuracy: 0.5770 - val_loss: 0.6901 - val_accuracy: 0.5503
Epoch 48/100
7000/7000 [=====] - 1s 139us/step - loss: 0.6741 - accuracy: 0.5843 - val_loss: 0.6875 - val_accuracy: 0.5543
Epoch 49/100
7000/7000 [=====] - 1s 142us/step - loss: 0.6743 - accuracy: 0.5806 - val_loss: 0.6865 - val_accuracy: 0.5527
Epoch 50/100
7000/7000 [=====] - 1s 140us/step - loss: 0.6751 - accuracy: 0.5726 - val_loss: 0.6872 - val_accuracy: 0.5507
Epoch 51/100
7000/7000 [=====] - 1s 133us/step - loss: 0.6737 - accuracy: 0.5859 - val_loss: 0.6936 - val_accuracy: 0.5427
Epoch 52/100
7000/7000 [=====] - 1s 132us/step - loss: 0.6744 - accuracy: 0.5841 - val_loss: 0.6850 - val_accuracy: 0.5600
Epoch 53/100
7000/7000 [=====] - 1s 132us/step - loss: 0.6740 - accuracy: 0.5784 - val_loss: 0.6847 - val_accuracy: 0.5557
Epoch 54/100
7000/7000 [=====] - 1s 132us/step - loss: 0.6739 - accuracy: 0.5840 - val_loss: 0.6859 - val_accuracy: 0.5530
Epoch 55/100
7000/7000 [=====] - 1s 132us/step - loss: 0.6725 - accuracy: 0.5841 - val_loss: 0.6859 - val_accuracy: 0.5480
Epoch 56/100
7000/7000 [=====] - 1s 132us/step - loss: 0.6732 - accuracy: 0.5773 - val_loss: 0.6844 - val_accuracy: 0.5483
Epoch 57/100
7000/7000 [=====] - 1s 137us/step - loss: 0.6727 - accuracy: 0.5857 - val_loss: 0.6845 - val_accuracy: 0.5530
Epoch 58/100
7000/7000 [=====] - 1s 157us/step - loss: 0.6730 - accuracy: 0.5847 - val_loss: 0.6842 - val_accuracy: 0.5527
Epoch 59/100
7000/7000 [=====] - 1s 132us/step - loss: 0.6739 - accuracy: 0.5829 - val_loss: 0.6843 - val_accuracy: 0.5527
Epoch 60/100
7000/7000 [=====] - 1s 132us/step - loss: 0.6724 - accuracy: 0.5850 - val_loss: 0.6834 - val_accuracy: 0.5623
Epoch 61/100
7000/7000 [=====] - 1s 147us/step - loss: 0.6717 - accuracy: 0.5830 - val_loss: 0.6864 - val_accuracy: 0.5540
Epoch 62/100
7000/7000 [=====] - 1s 147us/step - loss: 0.6735 - accuracy: 0.5837 - val_loss: 0.6862 - val_accuracy: 0.5590
Epoch 63/100
7000/7000 [=====] - 1s 141us/step - loss: 0.6738 - accuracy: 0.5833 - val_loss: 0.6837 - val_accuracy: 0.5650
Epoch 64/100
7000/7000 [=====] - 1s 137us/step - loss: 0.6731 - accuracy: 0.5821 - val_loss: 0.6845 - val_accuracy: 0.5543
Epoch 65/100
7000/7000 [=====] - 1s 133us/step - loss: 0.6715 - accuracy: 0.5849 - val_loss: 0.6872 - val_accuracy: 0.5627
Epoch 66/100
7000/7000 [=====] - 1s 133us/step - loss: 0.6729 - accuracy: 0.5860 - val_loss: 0.6876 - val_accuracy: 0.5557
Epoch 67/100
7000/7000 [=====] - 1s 165us/step - loss: 0.6721 - accuracy: 0.5827 - val_loss: 0.6841 - val_accuracy: 0.5643
Epoch 68/100
7000/7000 [=====] - 1s 138us/step - loss: 0.6720 - accuracy: 0.5836 - val_loss: 0.6845 - val_accuracy: 0.5657
Epoch 69/100

7000/7000 [=====] - 1s 133us/step - loss: 0.6699 - accuracy: 0.5904 - val_loss: 0.6834 - val_accuracy: 0.5600
Epoch 70/100
7000/7000 [=====] - 1s 133us/step - loss: 0.6716 - accuracy: 0.5839 - val_loss: 0.6832 - val_accuracy: 0.5663
Epoch 71/100
7000/7000 [=====] - 1s 133us/step - loss: 0.6711 - accuracy: 0.5853 - val_loss: 0.6834 - val_accuracy: 0.5713
Epoch 72/100
7000/7000 [=====] - 1s 134us/step - loss: 0.6726 - accuracy: 0.5840 - val_loss: 0.6843 - val_accuracy: 0.5540
Epoch 73/100
7000/7000 [=====] - 1s 135us/step - loss: 0.6722 - accuracy: 0.5784 - val_loss: 0.6837 - val_accuracy: 0.5620
Epoch 74/100
7000/7000 [=====] - 1s 142us/step - loss: 0.6717 - accuracy: 0.5789 - val_loss: 0.6839 - val_accuracy: 0.5617
Epoch 75/100
7000/7000 [=====] - 1s 146us/step - loss: 0.6708 - accuracy: 0.5853 - val_loss: 0.6844 - val_accuracy: 0.5617
Epoch 76/100
7000/7000 [=====] - 1s 137us/step - loss: 0.6698 - accuracy: 0.5923 - val_loss: 0.6840 - val_accuracy: 0.5570
Epoch 77/100
7000/7000 [=====] - 1s 144us/step - loss: 0.6692 - accuracy: 0.5899 - val_loss: 0.6844 - val_accuracy: 0.5540
Epoch 78/100
7000/7000 [=====] - 1s 144us/step - loss: 0.6690 - accuracy: 0.5864 - val_loss: 0.6890 - val_accuracy: 0.5557
Epoch 79/100
7000/7000 [=====] - 1s 138us/step - loss: 0.6714 - accuracy: 0.5791 - val_loss: 0.6860 - val_accuracy: 0.5563
Epoch 80/100
7000/7000 [=====] - 1s 133us/step - loss: 0.6679 - accuracy: 0.5906 - val_loss: 0.6867 - val_accuracy: 0.5600
Epoch 81/100
7000/7000 [=====] - 1s 141us/step - loss: 0.6698 - accuracy: 0.5880 - val_loss: 0.6871 - val_accuracy: 0.5520
Epoch 82/100
7000/7000 [=====] - 1s 142us/step - loss: 0.6684 - accuracy: 0.5913 - val_loss: 0.6860 - val_accuracy: 0.5620
Epoch 83/100
7000/7000 [=====] - 1s 140us/step - loss: 0.6690 - accuracy: 0.5884 - val_loss: 0.6875 - val_accuracy: 0.5567
Epoch 84/100
7000/7000 [=====] - 1s 135us/step - loss: 0.6712 - accuracy: 0.5907 - val_loss: 0.6880 - val_accuracy: 0.5567
Epoch 85/100
7000/7000 [=====] - 1s 141us/step - loss: 0.6694 - accuracy: 0.5870 - val_loss: 0.6971 - val_accuracy: 0.5477
Epoch 86/100
7000/7000 [=====] - 1s 136us/step - loss: 0.6693 - accuracy: 0.5873 - val_loss: 0.6861 - val_accuracy: 0.5550
Epoch 87/100
7000/7000 [=====] - 1s 138us/step - loss: 0.6659 - accuracy: 0.5973 - val_loss: 0.6880 - val_accuracy: 0.5553
Epoch 88/100
7000/7000 [=====] - 1s 141us/step - loss: 0.6657 - accuracy: 0.5954 - val_loss: 0.6883 - val_accuracy: 0.5597
Epoch 89/100
7000/7000 [=====] - 1s 145us/step - loss: 0.6674 - accuracy: 0.5881 - val_loss: 0.6902 - val_accuracy: 0.5480
Epoch 90/100
7000/7000 [=====] - 1s 148us/step - loss: 0.6672 - accuracy: 0.5879 - val_loss: 0.6916 - val_accuracy: 0.5490
Epoch 91/100
7000/7000 [=====] - 1s 144us/step - loss: 0.6671 - accuracy: 0.6000 - val_loss: 0.6885 - val_accuracy: 0.5507
Epoch 92/100
7000/7000 [=====] - 1s 137us/step - loss: 0.6658 - accuracy: 0.5950 - val_loss: 0.6898 - val_accuracy: 0.5527
Epoch 93/100
7000/7000 [=====] - 1s 134us/step - loss: 0.6659 - accuracy: 0.5999 - val_loss: 0.6894 - val_accuracy: 0.5610
Epoch 94/100
7000/7000 [=====] - 1s 133us/step - loss: 0.6650 - accuracy: 0.5994 - val_loss: 0.6896 - val_accuracy: 0.5507



```

Epoch 95/100
7000/7000 [=====] - 1s 137us/step - loss: 0.6683 - accuracy: 0.5906 - val_loss:
0.6886 - val_accuracy: 0.5477
Epoch 96/100
7000/7000 [=====] - 1s 135us/step - loss: 0.6669 - accuracy: 0.5896 - val_loss:
0.6872 - val_accuracy: 0.5580
Epoch 97/100
7000/7000 [=====] - 1s 136us/step - loss: 0.6651 - accuracy: 0.5943 - val_loss:
0.6882 - val_accuracy: 0.5620
Epoch 98/100
7000/7000 [=====] - 1s 133us/step - loss: 0.6650 - accuracy: 0.5936 - val_loss:
0.6895 - val_accuracy: 0.5550
Epoch 99/100
7000/7000 [=====] - 1s 142us/step - loss: 0.6647 - accuracy: 0.5967 - val_loss:
0.6902 - val_accuracy: 0.5597
Epoch 100/100
7000/7000 [=====] - 1s 134us/step - loss: 0.6663 - accuracy: 0.5910 - val_loss:
0.6935 - val_accuracy: 0.5537

```

In [54]:

```
plot_history(CNN_history)
```



In [56]:

```

y_train_pred = CNN_Model.predict(X_train.reshape(X_train.shape[0], X_train.shape[1], 1))
y_train_pred = np.where(y_train_pred < 0.5, 0, 1)
print(classification_report(y_train_pred, y_train))

yhat_xg_bestmodel = CNN_Model.predict(X_test.reshape(X_test.shape[0], X_test.shape[1], 1))
yhat_xg_bestmodel = np.where(yhat_xg_bestmodel < 0.5, 0, 1)
classi_reprt_xg = classification_report(yhat_xg_bestmodel, y_test)
print(classi_reprt_xg)

```

	precision	recall	f1-score	support
0	0.52	0.60	0.56	3006
1	0.66	0.58	0.62	3994
accuracy			0.59	7000
macro avg	0.59	0.59	0.59	7000
weighted avg	0.60	0.59	0.59	7000

	precision	recall	f1-score	support
0	0.48	0.57	0.52	1276
1	0.63	0.54	0.58	1724
accuracy			0.55	3000
macro avg	0.55	0.56	0.55	3000
weighted avg	0.57	0.55	0.56	3000

In []:

XGBoost - Binary - Word2Vec

In [18]:

```
# find the longest text to determine max_log_length
#length = [len(i) for i in X_encoded]
# max(length)
processed_data = pandas.read_csv("data_word2vec.csv")
processed_data
```

Out[18]:

	Unnamed: 0	0	1	2	3	4	5	6	7	8	...	42	43	44
0	0	0.046881	0.015105	0.092017	0.401055	0.090607	0.325169	0.297168	0.252889	0.431492	...	0.535171	0.110634	0.378137
1	1	0.060067	0.014853	0.000460	0.312015	0.175160	0.217285	0.240094	0.142320	0.265984	...	0.597190	0.062401	0.313474
2	2	0.061501	0.073079	0.033006	0.367333	0.133849	0.395722	0.239342	0.159746	0.462475	...	0.617383	0.151002	0.353643
3	3	0.657267	0.984994	0.582354	0.502777	0.039903	0.415720	0.229914	1.442688	1.223933	...	0.566871	0.438947	0.868708
4	4	0.248054	0.533198	0.372260	0.621143	0.119701	0.443986	0.274354	0.953549	0.972687	...	0.530449	0.079765	0.859161
...
9995	9995	0.020909	0.195280	0.117890	0.686506	0.043509	0.537159	0.368239	0.500929	0.594955	...	0.569722	0.234621	0.381084
9996	9996	0.072137	0.062902	0.131932	0.784530	0.031594	0.320527	0.411374	0.429918	0.767537	...	0.765428	0.371071	0.570000
9997	9997	0.102276	0.047583	0.443892	0.589184	0.012078	0.517465	0.362354	0.514404	0.876258	...	0.332553	0.242859	0.628040
9998	9998	0.006259	0.005005	0.135583	0.375453	0.426059	0.452148	0.376979	0.220052	0.507017	...	0.746278	0.164733	0.264629
9999	9999	0.541975	0.379249	0.034109	1.336829	0.046721	0.630619	0.939171	0.223359	0.905672	...	0.467232	0.915110	0.575144

10000 rows × 53 columns



In [19]:

```
# pad the data as each observation has a different length
#max_log_length = 2048 # larger than 2003
#X_processed = sequence.pad_sequences(X_encoded, maxlen=max_log_length)
```

In [20]:

```
#X_processed
y = processed_data[["Label"]].values
avg_data = processed_data.iloc[:, 1:-2].values
# y -= 1
avg_data
```

Out[20]:

```
array([[ -4.68814398e-02, -1.51045953e-02,  9.20168824e-02, ...,
        -3.73423553e-01, -4.77381344e-01,  3.07678892e-01],
       [-6.00674681e-02,  1.48525748e-02, -4.60325505e-04, ...,
        -2.95136365e-01, -3.68211810e-01,  3.59656972e-01],
       [-6.15012782e-02,  7.30789267e-02,  3.30062611e-02, ...,
        -2.75688761e-01, -4.71716609e-01,  4.11675429e-01],
       ...,
       [ 1.02276118e-01, -4.75825627e-02,  4.43892363e-01, ...,
        -4.47013392e-01, -7.74893663e-01, -1.60717556e-02],
       [ 6.25933642e-03, -5.00470836e-03,  1.35582868e-01, ...,
        -2.38587142e-01, -4.32617757e-01,  5.29272652e-01],
       [-5.41975147e-01,  3.79249226e-01, -3.41089983e-02, ...,
        -4.11995113e-01, -1.29827176e-01,  3.93914914e-01]])
```

In [21]:

```
avg_data = np.array(avg_data)
y = y.reshape((y.shape[0], 1))
```

In [22]:

```
# split 70% train and 30% test data
X_train, X_test, y_train, y_test = train_test_split(avg_data, y, test_size=0.30, random_state=0)
y_train_cat = to_categorical(y_train)
y_test_cat = to_categorical(y_test)
```

In []:

In [27]:

```
import xgboost as xgb
```

In [23]:

```
xg_rdsearch = xgb.XGBClassifier(alpha = 0.1)
```

In [24]:

```
param_rdsearch_xg = {'n_estimators':[100, 150, 200, 250, 300, 350, 400, 450, 500,550, 600, 650, 700, 750,
                                     850, 900, 950, 1000],
                     'learning_rate':[0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1, 1.1, 1.2, 1.3, 1.4,
                     'max_depth':[1,2],
                     'gamma': [0, 0.25, 0.5, 0.75, 1, 1.25, 1.5, 1.75, 2, 2.25, 2.5, 2.75, 3, 3.25, 3.5, 3
                               4.5, 4.75, 5]}
```

```
xg_rdsearch = RandomizedSearchCV(xg_rdsearch, param_rdsearch_xg, cv = 5, scoring = 'roc_auc',
                                refit = True, n_jobs=-1, verbose = 5)
```

In [25]:

```
xg_rdsearch.fit(X_train, y_train)
```

Fitting 5 folds for each of 10 candidates, totalling 50 fits

C:\Users\14264\Anaconda3\envs\tf_gpu\lib\site-packages\xgboost\sklearn.py:888: UserWarning: The use of label encoder in XGBClassifier is deprecated and will be removed in a future release. To remove this warning, do the following: 1) Pass option use_label_encoder=False when constructing XGBClassifier object; and 2) Encode your labels (y) as integers starting with 0, i.e. 0, 1, 2, ..., [num_class - 1].

warnings.warn(label_encoder_deprecation_msg, UserWarning)

C:\Users\14264\Anaconda3\envs\tf_gpu\lib\site-packages\sklearn\utils\validation.py:63:

DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

return f(*args, **kwargs)

[01:25:37] WARNING: ..\src\learner.cc:1061: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

Out[25]:

```
RandomizedSearchCV(cv=5,
                  estimator=XGBClassifier(alpha=0.1, base_score=None,
                                          booster=None, colsample_bylevel=None,
                                          colsample_bynode=None,
                                          colsample_bytree=None, gamma=None,
                                          gpu_id=None, importance_type='gain',
                                          interaction_constraints=None,
                                          learning_rate=None,
                                          max_delta_step=None, max_depth=None,
                                          min_child_weight=None, missing=nan,
                                          monotone_constraints=None,
                                          n_estim...,
                                          verbosity=None),
                  n_jobs=-1,
                  param_distributions={'gamma': [0, 0.25, 0.5, 0.75, 1, 1.25,
                                                1.5, 1.75, 2, 2.25, 2.5, 2.75,
                                                3, 3.25, 3.5, 3.75, 4, 4.25,
                                                4.5, 4.75, 5],
                                     'learning_rate': [0.1, 0.2, 0.3, 0.4,
                                                       0.5, 0.6, 0.7, 0.8,
                                                       0.9, 1, 1.1, 1.2, 1.3,
                                                       1.4, 1.5, 1.6],
                                     'max_depth': [1, 2],
                                     'n_estimators': [100, 150, 200, 250,
                                                       300, 350, 400, 450,
                                                       500, 550, 600, 650,
                                                       700, 750, 800, 850,
                                                       900, 950, 1000]},
                  scoring='roc_auc', verbose=5)
```

In [26]:

```
xg_bestparams = xg_rdsearch.best_params_
xg_bestparams
```

Out[26]:

```
{'n_estimators': 650, 'max_depth': 1, 'learning_rate': 0.3, 'gamma': 4.25}
```

In [27]:

```
xg_bestmodel = xg_rdsearch.best_estimator_
xg_bestmodel
```

Out[27]:

```
XGBClassifier(alpha=0.1, base_score=0.5, booster='gbtree', colsample_bylevel=1,
              colsample_bynode=1, colsample_bytree=1, gamma=4.25, gpu_id=-1,
              importance_type='gain', interaction_constraints='',
              learning_rate=0.3, max_delta_step=0, max_depth=1,
              min_child_weight=1, missing=nan, monotone_constraints='()',
              n_estimators=650, n_jobs=8, num_parallel_tree=1, random_state=0,
              reg_alpha=0.100000001, reg_lambda=1, scale_pos_weight=1,
              subsample=1, tree_method='exact', validate_parameters=1,
              verbosity=None)
```

In [78]:

```
y_train_pred = xg_bestmodel.predict(X_train)
```

In [79]:

```
print(classification_report(y_train_pred,y_train))
```

	precision	recall	f1-score	support
0	0.74	0.75	0.75	3467
1	0.75	0.74	0.75	3533
accuracy			0.75	7000
macro avg	0.75	0.75	0.75	7000
weighted avg	0.75	0.75	0.75	7000

In [28]:

```
yhat_xg_bestmodel = xg_bestmodel.predict(X_test)
```

In [29]:

```
classi_reprt_xg = classification_report(y_test, yhat_xg_bestmodel)
```

In [30]:

```
print(classi_reprt_xg)
```

	precision	recall	f1-score	support
0	0.72	0.70	0.71	1494
1	0.71	0.73	0.72	1506
accuracy			0.72	3000
macro avg	0.72	0.72	0.72	3000
weighted avg	0.72	0.72	0.72	3000

CNN - Binary - word2vec

In [25]:

```
CNN_Model = Sequential()
CNN_Model.add(Conv1D(10, 20, activation = 'sigmoid', input_shape = (50, 1)))
CNN_Model.add(Conv1D(20, 8, activation = 'sigmoid'))
CNN_Model.add(Conv1D(30, 5, activation = 'sigmoid'))
CNN_Model.add(GlobalMaxPooling1D())
CNN_Model.add(Dense(10, activation = 'sigmoid'))
CNN_Model.add(Dense(1, activation = 'sigmoid'))
CNN_Model.compile(optimizer = 'adam', loss = 'binary_crossentropy', metrics = ['accuracy'])
CNN_Model.summary()
```

Model: "sequential_3"

Layer (type)	Output Shape	Param #
conv1d_6 (Conv1D)	(None, 31, 10)	210
conv1d_7 (Conv1D)	(None, 24, 20)	1620
conv1d_8 (Conv1D)	(None, 20, 30)	3030
global_max_pooling1d_2 (Glob	(None, 30)	0
dense_4 (Dense)	(None, 10)	310
dense_5 (Dense)	(None, 1)	11
Total params: 5,181		
Trainable params: 5,181		
Non-trainable params: 0		

In [26]:

```
CNN_history = CNN_Model.fit(X_train.reshape((X_train.shape[0], X_train.shape[1], 1)), y_train,
                             epochs = 100, batch_size = 50,
                             validation_data = (X_test.reshape((X_test.shape[0], X_test.shape[1], 1)), y_test))
```

Train on 7000 samples, validate on 3000 samples

Epoch 1/100

7000/7000 [=====] - 1s 94us/step - loss: 0.6972 - accuracy: 0.4926 - val_loss: 0.6931 - val_accuracy: 0.5020

Epoch 2/100

7000/7000 [=====] - 0s 62us/step - loss: 0.6932 - accuracy: 0.4996 - val_loss:

```
7000/7000 [-----] - 0s 62us/step - loss: 0.6932 - accuracy: 0.4990 - val_loss:
0.6932 - val_accuracy: 0.4980
Epoch 3/100
7000/7000 [=====] - 0s 61us/step - loss: 0.6934 - accuracy: 0.5023 - val_loss:
0.6928 - val_accuracy: 0.5020
Epoch 4/100
7000/7000 [=====] - 0s 61us/step - loss: 0.6932 - accuracy: 0.5120 - val_loss:
0.6925 - val_accuracy: 0.4980
Epoch 5/100
7000/7000 [=====] - 0s 60us/step - loss: 0.6927 - accuracy: 0.5126 - val_loss:
0.6918 - val_accuracy: 0.4980
Epoch 6/100
7000/7000 [=====] - 0s 68us/step - loss: 0.6897 - accuracy: 0.5503 - val_loss:
0.6862 - val_accuracy: 0.5677
Epoch 7/100
7000/7000 [=====] - 0s 61us/step - loss: 0.6776 - accuracy: 0.5986 - val_loss:
0.6658 - val_accuracy: 0.6140
Epoch 8/100
7000/7000 [=====] - 0s 61us/step - loss: 0.6586 - accuracy: 0.6176 - val_loss:
0.6487 - val_accuracy: 0.6333
Epoch 9/100
7000/7000 [=====] - 0s 61us/step - loss: 0.6458 - accuracy: 0.6296 - val_loss:
0.6398 - val_accuracy: 0.6440
Epoch 10/100
7000/7000 [=====] - 0s 64us/step - loss: 0.6374 - accuracy: 0.6390 - val_loss:
0.6306 - val_accuracy: 0.6533
Epoch 11/100
7000/7000 [=====] - 0s 61us/step - loss: 0.6286 - accuracy: 0.6490 - val_loss:
0.6290 - val_accuracy: 0.6523
Epoch 12/100
7000/7000 [=====] - 0s 61us/step - loss: 0.6187 - accuracy: 0.6586 - val_loss:
0.6091 - val_accuracy: 0.6690
Epoch 13/100
7000/7000 [=====] - 0s 61us/step - loss: 0.6117 - accuracy: 0.6693 - val_loss:
0.5999 - val_accuracy: 0.6780
Epoch 14/100
7000/7000 [=====] - 0s 60us/step - loss: 0.6038 - accuracy: 0.6759 - val_loss:
0.5930 - val_accuracy: 0.6843
Epoch 15/100
7000/7000 [=====] - 0s 61us/step - loss: 0.5992 - accuracy: 0.6794 - val_loss:
0.5879 - val_accuracy: 0.6923
Epoch 16/100
7000/7000 [=====] - 0s 60us/step - loss: 0.5976 - accuracy: 0.6817 - val_loss:
0.5848 - val_accuracy: 0.6930
Epoch 17/100
7000/7000 [=====] - 0s 60us/step - loss: 0.5925 - accuracy: 0.6800 - val_loss:
0.5826 - val_accuracy: 0.6973
Epoch 18/100
7000/7000 [=====] - 0s 61us/step - loss: 0.5906 - accuracy: 0.6907 - val_loss:
0.5811 - val_accuracy: 0.6953
Epoch 19/100
7000/7000 [=====] - 0s 60us/step - loss: 0.5866 - accuracy: 0.6917 - val_loss:
0.5817 - val_accuracy: 0.6920
Epoch 20/100
7000/7000 [=====] - 0s 61us/step - loss: 0.5850 - accuracy: 0.6940 - val_loss:
0.5745 - val_accuracy: 0.7020
Epoch 21/100
7000/7000 [=====] - 0s 65us/step - loss: 0.5827 - accuracy: 0.6936 - val_loss:
0.5740 - val_accuracy: 0.7003
Epoch 22/100
7000/7000 [=====] - 0s 65us/step - loss: 0.5790 - accuracy: 0.7017 - val_loss:
0.5746 - val_accuracy: 0.6993
Epoch 23/100
7000/7000 [=====] - 0s 60us/step - loss: 0.5763 - accuracy: 0.7009 - val_loss:
0.5720 - val_accuracy: 0.7007
Epoch 24/100
7000/7000 [=====] - 0s 66us/step - loss: 0.5769 - accuracy: 0.7010 - val_loss:
0.5692 - val_accuracy: 0.7017
Epoch 25/100
7000/7000 [=====] - 0s 66us/step - loss: 0.5749 - accuracy: 0.7020 - val_loss:
0.5660 - val_accuracy: 0.7100
Epoch 26/100
7000/7000 [=====] - 0s 63us/step - loss: 0.5726 - accuracy: 0.7060 - val_loss:
0.5664 - val_accuracy: 0.7070
Epoch 27/100
7000/7000 [=====] - 0s 61us/step - loss: 0.5697 - accuracy: 0.7110 - val_loss:
0.5663 - val_accuracy: 0.7087
Epoch 28/100
```

Epoch 28/100
7000/7000 [=====] - 0s 60us/step - loss: 0.5688 - accuracy: 0.7086 - val_loss: 0.5626 - val_accuracy: 0.7120
Epoch 29/100
7000/7000 [=====] - 0s 69us/step - loss: 0.5670 - accuracy: 0.7073 - val_loss: 0.5613 - val_accuracy: 0.7133
Epoch 30/100
7000/7000 [=====] - 0s 61us/step - loss: 0.5662 - accuracy: 0.7107 - val_loss: 0.5598 - val_accuracy: 0.7120
Epoch 31/100
7000/7000 [=====] - 0s 65us/step - loss: 0.5642 - accuracy: 0.7094 - val_loss: 0.5596 - val_accuracy: 0.7130
Epoch 32/100
7000/7000 [=====] - 0s 60us/step - loss: 0.5643 - accuracy: 0.7111 - val_loss: 0.5617 - val_accuracy: 0.7060
Epoch 33/100
7000/7000 [=====] - 0s 60us/step - loss: 0.5626 - accuracy: 0.7134 - val_loss: 0.5591 - val_accuracy: 0.7100
Epoch 34/100
7000/7000 [=====] - 0s 61us/step - loss: 0.5613 - accuracy: 0.7164 - val_loss: 0.5574 - val_accuracy: 0.7127
Epoch 35/100
7000/7000 [=====] - 0s 62us/step - loss: 0.5598 - accuracy: 0.7156 - val_loss: 0.5562 - val_accuracy: 0.7143
Epoch 36/100
7000/7000 [=====] - 0s 63us/step - loss: 0.5614 - accuracy: 0.7177 - val_loss: 0.5592 - val_accuracy: 0.7100
Epoch 37/100
7000/7000 [=====] - 0s 60us/step - loss: 0.5609 - accuracy: 0.7124 - val_loss: 0.5550 - val_accuracy: 0.7107
Epoch 38/100
7000/7000 [=====] - 0s 61us/step - loss: 0.5600 - accuracy: 0.7199 - val_loss: 0.5588 - val_accuracy: 0.7093
Epoch 39/100
7000/7000 [=====] - 0s 63us/step - loss: 0.5581 - accuracy: 0.7186 - val_loss: 0.5547 - val_accuracy: 0.7117
Epoch 40/100
7000/7000 [=====] - 0s 65us/step - loss: 0.5571 - accuracy: 0.7186 - val_loss: 0.5557 - val_accuracy: 0.7100
Epoch 41/100
7000/7000 [=====] - 0s 67us/step - loss: 0.5581 - accuracy: 0.7139 - val_loss: 0.5623 - val_accuracy: 0.7060
Epoch 42/100
7000/7000 [=====] - 0s 63us/step - loss: 0.5593 - accuracy: 0.7173 - val_loss: 0.5563 - val_accuracy: 0.7103
Epoch 43/100
7000/7000 [=====] - 0s 63us/step - loss: 0.5542 - accuracy: 0.7201 - val_loss: 0.5536 - val_accuracy: 0.7130
Epoch 44/100
7000/7000 [=====] - 0s 64us/step - loss: 0.5571 - accuracy: 0.7160 - val_loss: 0.5525 - val_accuracy: 0.7150
Epoch 45/100
7000/7000 [=====] - 0s 66us/step - loss: 0.5546 - accuracy: 0.7194 - val_loss: 0.5523 - val_accuracy: 0.7117
Epoch 46/100
7000/7000 [=====] - 0s 66us/step - loss: 0.5524 - accuracy: 0.7189 - val_loss: 0.5509 - val_accuracy: 0.7103
Epoch 47/100
7000/7000 [=====] - 0s 67us/step - loss: 0.5529 - accuracy: 0.7229 - val_loss: 0.5525 - val_accuracy: 0.7130
Epoch 48/100
7000/7000 [=====] - 0s 70us/step - loss: 0.5501 - accuracy: 0.7206 - val_loss: 0.5625 - val_accuracy: 0.7073
Epoch 49/100
7000/7000 [=====] - 0s 64us/step - loss: 0.5523 - accuracy: 0.7217 - val_loss: 0.5532 - val_accuracy: 0.7140
Epoch 50/100
7000/7000 [=====] - 0s 64us/step - loss: 0.5490 - accuracy: 0.7266 - val_loss: 0.5481 - val_accuracy: 0.7143
Epoch 51/100
7000/7000 [=====] - 1s 73us/step - loss: 0.5497 - accuracy: 0.7260 - val_loss: 0.5592 - val_accuracy: 0.7050
Epoch 52/100
7000/7000 [=====] - 0s 69us/step - loss: 0.5519 - accuracy: 0.7236 - val_loss: 0.5474 - val_accuracy: 0.7150
Epoch 53/100
7000/7000 [=====] - 0s 65us/step - loss: 0.5489 - accuracy: 0.7213 - val_loss: 0.5528 - val_accuracy: 0.7102

```
0.5529 - val_accuracy: 0.7103
Epoch 54/100
7000/7000 [=====] - 0s 65us/step - loss: 0.5470 - accuracy: 0.7293 - val_loss:
0.5472 - val_accuracy: 0.7137
Epoch 55/100
7000/7000 [=====] - 0s 62us/step - loss: 0.5476 - accuracy: 0.7290 - val_loss:
0.5460 - val_accuracy: 0.7147
Epoch 56/100
7000/7000 [=====] - 0s 63us/step - loss: 0.5481 - accuracy: 0.7219 - val_loss:
0.5456 - val_accuracy: 0.7173
Epoch 57/100
7000/7000 [=====] - 0s 61us/step - loss: 0.5467 - accuracy: 0.7270 - val_loss:
0.5461 - val_accuracy: 0.7140
Epoch 58/100
7000/7000 [=====] - 0s 61us/step - loss: 0.5464 - accuracy: 0.7271 - val_loss:
0.5480 - val_accuracy: 0.7153
Epoch 59/100
7000/7000 [=====] - 0s 63us/step - loss: 0.5433 - accuracy: 0.7297 - val_loss:
0.5447 - val_accuracy: 0.7150
Epoch 60/100
7000/7000 [=====] - 0s 63us/step - loss: 0.5433 - accuracy: 0.7271 - val_loss:
0.5461 - val_accuracy: 0.7150
Epoch 61/100
7000/7000 [=====] - 0s 63us/step - loss: 0.5435 - accuracy: 0.7307 - val_loss:
0.5549 - val_accuracy: 0.7063
Epoch 62/100
7000/7000 [=====] - 0s 66us/step - loss: 0.5425 - accuracy: 0.7277 - val_loss:
0.5465 - val_accuracy: 0.7147
Epoch 63/100
7000/7000 [=====] - 0s 69us/step - loss: 0.5419 - accuracy: 0.7329 - val_loss:
0.5448 - val_accuracy: 0.7187
Epoch 64/100
7000/7000 [=====] - 0s 67us/step - loss: 0.5431 - accuracy: 0.7284 - val_loss:
0.5451 - val_accuracy: 0.7180
Epoch 65/100
7000/7000 [=====] - 0s 62us/step - loss: 0.5402 - accuracy: 0.7287 - val_loss:
0.5432 - val_accuracy: 0.7177
Epoch 66/100
7000/7000 [=====] - 0s 62us/step - loss: 0.5418 - accuracy: 0.7311 - val_loss:
0.5459 - val_accuracy: 0.7150
Epoch 67/100
7000/7000 [=====] - 0s 62us/step - loss: 0.5414 - accuracy: 0.7297 - val_loss:
0.5436 - val_accuracy: 0.7177
Epoch 68/100
7000/7000 [=====] - 0s 64us/step - loss: 0.5394 - accuracy: 0.7320 - val_loss:
0.5434 - val_accuracy: 0.7190
Epoch 69/100
7000/7000 [=====] - 0s 63us/step - loss: 0.5406 - accuracy: 0.7296 - val_loss:
0.5436 - val_accuracy: 0.7163
Epoch 70/100
7000/7000 [=====] - 0s 64us/step - loss: 0.5397 - accuracy: 0.7356 - val_loss:
0.5411 - val_accuracy: 0.7210
Epoch 71/100
7000/7000 [=====] - 1s 75us/step - loss: 0.5389 - accuracy: 0.7289 - val_loss:
0.5438 - val_accuracy: 0.7203
Epoch 72/100
7000/7000 [=====] - 0s 62us/step - loss: 0.5410 - accuracy: 0.7296 - val_loss:
0.5491 - val_accuracy: 0.7103
Epoch 73/100
7000/7000 [=====] - 0s 67us/step - loss: 0.5354 - accuracy: 0.7390 - val_loss:
0.5416 - val_accuracy: 0.7173
Epoch 74/100
7000/7000 [=====] - 0s 64us/step - loss: 0.5360 - accuracy: 0.7346 - val_loss:
0.5418 - val_accuracy: 0.7170
Epoch 75/100
7000/7000 [=====] - 0s 63us/step - loss: 0.5347 - accuracy: 0.7366 - val_loss:
0.5421 - val_accuracy: 0.7210
Epoch 76/100
7000/7000 [=====] - 0s 63us/step - loss: 0.5349 - accuracy: 0.7356 - val_loss:
0.5420 - val_accuracy: 0.7213
Epoch 77/100
7000/7000 [=====] - 0s 62us/step - loss: 0.5361 - accuracy: 0.7327 - val_loss:
0.5399 - val_accuracy: 0.7227
Epoch 78/100
7000/7000 [=====] - 0s 61us/step - loss: 0.5362 - accuracy: 0.7331 - val_loss:
0.5402 - val_accuracy: 0.7173
Epoch 79/100
7000/7000 [=====] - 0s 62us/step - loss: 0.5338 - accuracy: 0.7361 - val_loss:
0.5402 - val_accuracy: 0.7173
```




```

/0000/7000 [=====] - 0s 62us/step - loss: 0.5332 - accuracy: 0.7361 - val_loss:
0.5384 - val_accuracy: 0.7210
Epoch 80/100
7000/7000 [=====] - 0s 64us/step - loss: 0.5322 - accuracy: 0.7401 - val_loss:
0.5654 - val_accuracy: 0.7040
Epoch 81/100
7000/7000 [=====] - 0s 64us/step - loss: 0.5320 - accuracy: 0.7411 - val_loss:
0.5415 - val_accuracy: 0.7147
Epoch 82/100
7000/7000 [=====] - 0s 62us/step - loss: 0.5341 - accuracy: 0.7333 - val_loss:
0.5556 - val_accuracy: 0.7143
Epoch 83/100
7000/7000 [=====] - 0s 63us/step - loss: 0.5339 - accuracy: 0.7314 - val_loss:
0.5467 - val_accuracy: 0.7180
Epoch 84/100
7000/7000 [=====] - 0s 63us/step - loss: 0.5304 - accuracy: 0.7401 - val_loss:
0.5384 - val_accuracy: 0.7240
Epoch 85/100
7000/7000 [=====] - 0s 63us/step - loss: 0.5292 - accuracy: 0.7427 - val_loss:
0.5436 - val_accuracy: 0.7180
Epoch 86/100
7000/7000 [=====] - 0s 63us/step - loss: 0.5296 - accuracy: 0.7404 - val_loss:
0.5375 - val_accuracy: 0.7267
Epoch 87/100
7000/7000 [=====] - 0s 62us/step - loss: 0.5315 - accuracy: 0.7363 - val_loss:
0.5385 - val_accuracy: 0.7197
Epoch 88/100
7000/7000 [=====] - 0s 63us/step - loss: 0.5290 - accuracy: 0.7386 - val_loss:
0.5382 - val_accuracy: 0.7190
Epoch 89/100
7000/7000 [=====] - 0s 62us/step - loss: 0.5290 - accuracy: 0.7414 - val_loss:
0.5384 - val_accuracy: 0.7223
Epoch 90/100
7000/7000 [=====] - 0s 64us/step - loss: 0.5286 - accuracy: 0.7404 - val_loss:
0.5362 - val_accuracy: 0.7263
Epoch 91/100
7000/7000 [=====] - 0s 63us/step - loss: 0.5281 - accuracy: 0.7400 - val_loss:
0.5341 - val_accuracy: 0.7227
Epoch 92/100
7000/7000 [=====] - 0s 64us/step - loss: 0.5277 - accuracy: 0.7376 - val_loss:
0.5351 - val_accuracy: 0.7197
Epoch 93/100
7000/7000 [=====] - 0s 64us/step - loss: 0.5248 - accuracy: 0.7430 - val_loss:
0.5362 - val_accuracy: 0.7180
Epoch 94/100
7000/7000 [=====] - 0s 63us/step - loss: 0.5271 - accuracy: 0.7389 - val_loss:
0.5477 - val_accuracy: 0.7170
Epoch 95/100
7000/7000 [=====] - 0s 62us/step - loss: 0.5255 - accuracy: 0.7417 - val_loss:
0.5340 - val_accuracy: 0.7310
Epoch 96/100
7000/7000 [=====] - 0s 67us/step - loss: 0.5242 - accuracy: 0.7421 - val_loss:
0.5333 - val_accuracy: 0.7237
Epoch 97/100
7000/7000 [=====] - 0s 64us/step - loss: 0.5225 - accuracy: 0.7411 - val_loss:
0.5377 - val_accuracy: 0.7203
Epoch 98/100
7000/7000 [=====] - 0s 63us/step - loss: 0.5245 - accuracy: 0.7399 - val_loss:
0.5372 - val_accuracy: 0.7207
Epoch 99/100
7000/7000 [=====] - 0s 63us/step - loss: 0.5232 - accuracy: 0.7394 - val_loss:
0.5342 - val_accuracy: 0.7270
Epoch 100/100
7000/7000 [=====] - 0s 63us/step - loss: 0.5227 - accuracy: 0.7433 - val_loss:
0.5357 - val_accuracy: 0.7217

```

In [19]:

```
plot_history(CNN_history)
```



In [32]:

```
y_train_pred = CNN_Model.predict(X_train.reshape(X_train.shape[0], X_train.shape[1], 1))
y_train_pred = np.where(y_train_pred < 0.5, 0, 1)
print(classification_report(y_train_pred, y_train))

yhat_xg_bestmodel = CNN_Model.predict(X_test.reshape(X_test.shape[0], X_test.shape[1], 1))
yhat_xg_bestmodel = np.where(yhat_xg_bestmodel < 0.5, 0, 1)
classi_reprt_xg = classification_report(yhat_xg_bestmodel, y_test)
print(classi_reprt_xg)
```

	precision	recall	f1-score	support
0	0.81	0.71	0.76	3957
1	0.68	0.78	0.72	3043
accuracy			0.74	7000
macro avg	0.74	0.75	0.74	7000
weighted avg	0.75	0.74	0.74	7000

	precision	recall	f1-score	support
0	0.79	0.69	0.74	1711
1	0.65	0.76	0.70	1289
accuracy			0.72	3000
macro avg	0.72	0.73	0.72	3000
weighted avg	0.73	0.72	0.72	3000

In []: